

# Surface Detection by Robot Movements

Marian Dumitrascu

March 19, 2019

## Introduction

For this project I choosed a Kaggle.com open competition project. This is *CareerCon 2019 - Help Navigate Robots*. Here is the description of the project from Kaggle:

*In this competition, you'll help robots recognize the floor surface they're standing on using data collected from Inertial Measurement Units (IMU sensors). We've collected IMU sensor data while driving a small mobile robot over different floor surfaces on the university premises. The task is to predict which one of the nine floor types (carpet, tiles, concrete) the robot is on using sensor data such as acceleration and velocity. Succeed and you'll help improve the navigation of robots without assistance across many different surfaces, so they won't fall down on the job.*

The task is challenging, but I believe I can obtain a decent result using techniques and tools learned in this course. Please note that many code chunks are not showing in the PDF for aesthetic reasons. You can find the code behind many of the graphs and images in the Rmd version.

## Report Structure

1. First I will describe data provided and the output expected
2. Then, will perform data analysis, visualization and get insights
3. pre-process and transform data
4. decide the model to use, measure its performance and tune it
5. perform the final data prediction and show the results, submit to Kaggle
6. draw some conclusions

## Report and Data Location

I also keep this project on GitHub here: [https://github.com/mariandumitrascu/ph125\\_9\\_HelpRobotsNavigate](https://github.com/mariandumitrascu/ph125_9_HelpRobotsNavigate)  
Data loaded by the R scripts is kept on an AWS public S3 bucket, to be easily loaded. This will be available for the duration of grading.

## Data Description

Input data from Kaggle consists in 4 files:

- *X\_train.csv* and *X\_test.csv* - the input data, covering 10 sensor channels and 128 measurements per time series plus three ID columns:
  - *row\_id*: The ID for this row.
  - *series\_id*: ID number for the measurement series. Foreign key to *y\_train/sample\_submission*.
  - *measurement\_number*: Measurement number within the series.

The orientation channels encode the current angles of how the robot is oriented as a quaternion (see Wikipedia). Angular velocity describes the angle and speed of motion, and linear acceleration components describe how the speed is changing at different times. The 10 sensor channels are:

- *orientation\_X*
  - *orientation\_Y*
  - *orientation\_Z*
  - *orientation\_W*
  - *angular\_velocity\_X*
  - *angular\_velocity\_Y*
  - *angular\_velocity\_Z*
  - *linear\_acceleration\_X*
  - *linear\_acceleration\_Y*
  - *linear\_acceleration\_Z*
- *y\_train.csv* - the surfaces for training set.
    - *series\_id*: ID number for the measurement series.
    - *group\_id*: ID number for all of the measurements taken in a recording session. Provided for the training set only, to enable more cross validation strategies.
    - *surface*: labels or classes of the training data. this is the element that need to be predicted
  - *sample\_submission.csv* - a sample submission file in the correct format.

In this report I will split the training data into two partitions, will fit a model on the first one, and measure it's accuracy on the second. I will also use a small part of data to make it run faster. The R script for generating the final results will use all data.

## Data Analysis

I will make the following assumptions about observations:

- all observations are made using the same robot
- the interval between the 128 observations for each seeries is always the same.
- the surface is a plane, no stairs, hills or valleys

From a physycst perspective there are thre forces involved: gravitation force, robot propulsion force, and friction force. Gravitation force is constant. Friction force depends on the surface by a coefficient and propulsion is an unknown variable. We need to basically determine the friction coefficient based on a movement pattern. Moving objects will travel longer if the surface has a lower friction than on a surface with higher friction. On the other side, changing direction can be teeper on a surface with higher friction.

## First Insights

Quick look at the first 5 rows of the training data:

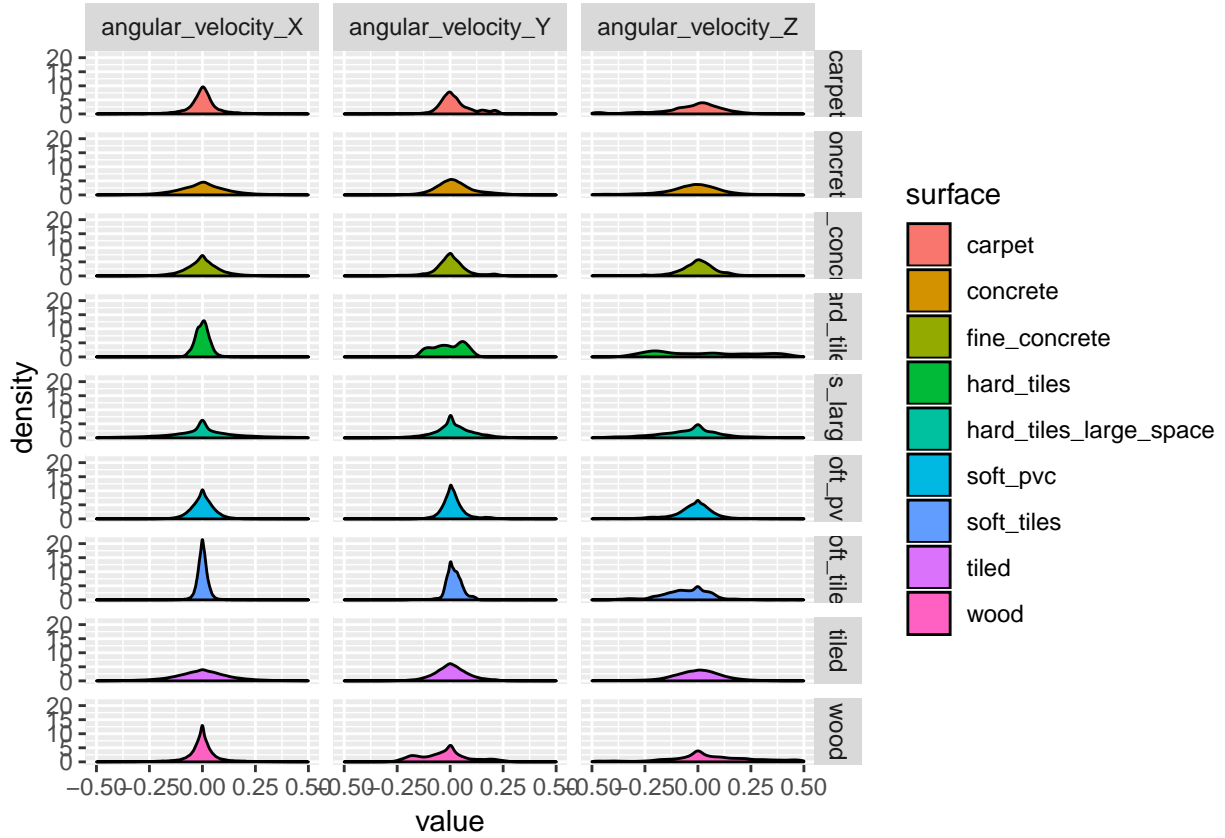
row_id	series_id	measurement_number	orientation_X	orientation_Y	orientation_Z	orientation_W	angular_
0_0	0	0	-0.75853	-0.63435	-0.10488	-0.10597	
0_1	0	1	-0.75853	-0.63434	-0.10490	-0.10600	
0_2	0	2	-0.75853	-0.63435	-0.10492	-0.10597	
0_3	0	3	-0.75852	-0.63436	-0.10495	-0.10597	
0_4	0	4	-0.75852	-0.63435	-0.10495	-0.10596	

## Data Distribution

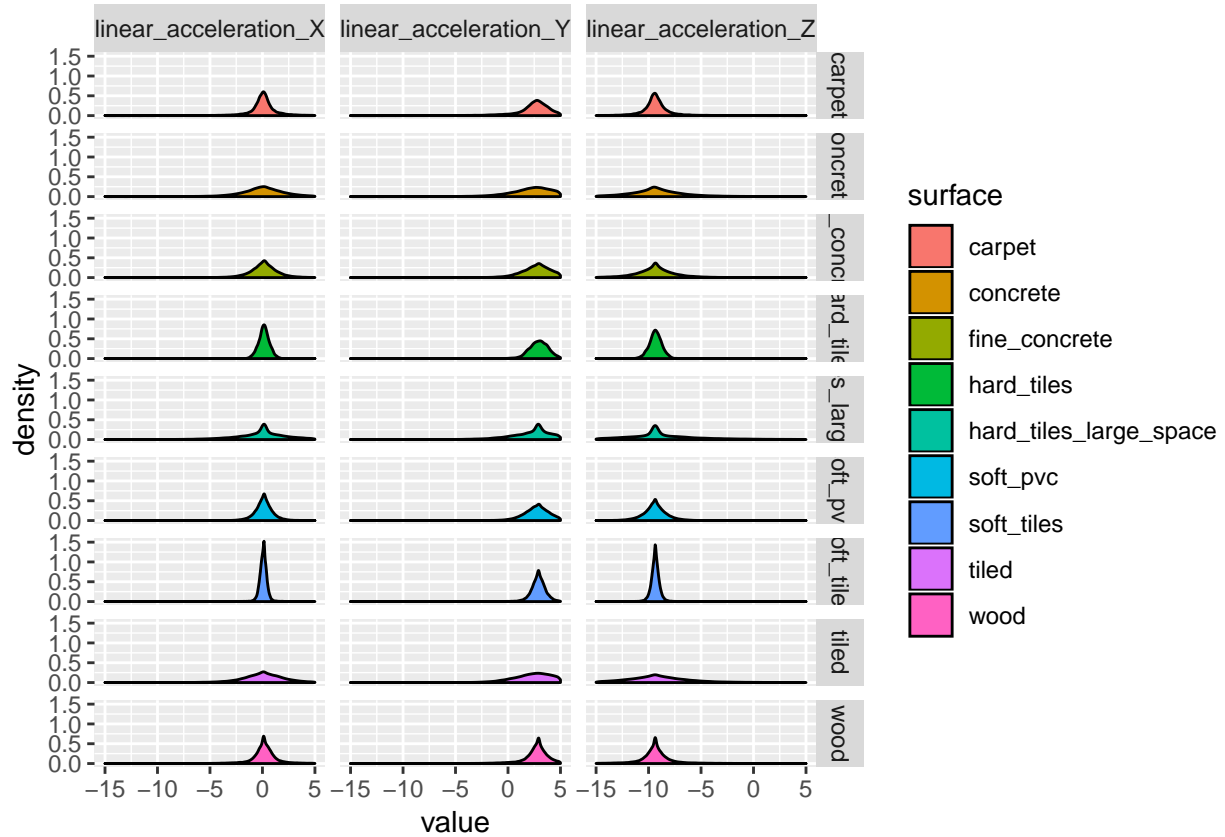
Here is a distributon of measurements by surface in the training set:

surface	measurements
carpet	189
concrete	779
fine_concrete	363
hard_tiles	21
hard_tiles_large_space	308
soft_pvc	732
soft_tiles	297
tiled	514
wood	607

Angular velocity data is produced by a magnetostatic sensor, it indicates the angular speed the robot is movig in reference with earth orientation. Here is the distribution of angular velocity by surface:



Linear acceleration data is produced by an inertial sensor. We can approximate this later with linear distacne if we consider the unit of time to be 1. Here is the distribution of linear acceleration by surface:



We can observe noticeable differences in distribution of these variables by surface.

Orientation data comes from a gyroscope sensor. To draw the distribution of orientation, we'll convert quaternion values to euler angles which are more intuitive and easier to interpret. Euler angles provide a way to represent the 3D orientation of an object using a combination of three rotations about different axes:

- roll - rotation around x, noted *phi*,
- pitch - rotation around y, noted *theta*,
- yaw - rotation around z, noted *psi*

See the image below (from: <http://www.chrobotics.com/library/understanding-euler-angles>)




Figure 1: Euler Angles

To convert quaternions to euler angles I used `Q2EA` function from `orientlib` package. (See: <https://www.rdocumentation.org/packages/orientlib/versions/0.10.3>)

```
# define a function to convert quaternion values to euler angles.
convert_quaternions_to_euler <- function(a_dataset){

  # use Q2EA from RSpincalc to convert quaternions to euler angles
  Q <- a_dataset %>% select(
    orientation_X,
    orientation_Y,
    orientation_Z,
    orientation_W) %>%
    as.matrix()

  euler_matrix <- Q2EA(Q,
    EulerOrder='xyz',
    tol = 10 * .Machine$double.eps,
    ichk = FALSE,
    ignoreAllChk = FALSE)

  # add the new columns to the dataset
  a_dataset <- a_dataset %>% mutate(
    phi = euler_matrix[,1],
    theta = euler_matrix[,2],
    psi = euler_matrix[,3])

  # remove quaternion columns
  a_dataset <- a_dataset %>% select(
    -orientation_X,
    -orientation_Y,
    -orientation_Z,
    -orientation_W)

  # return the new dataset
  a_dataset
}
```

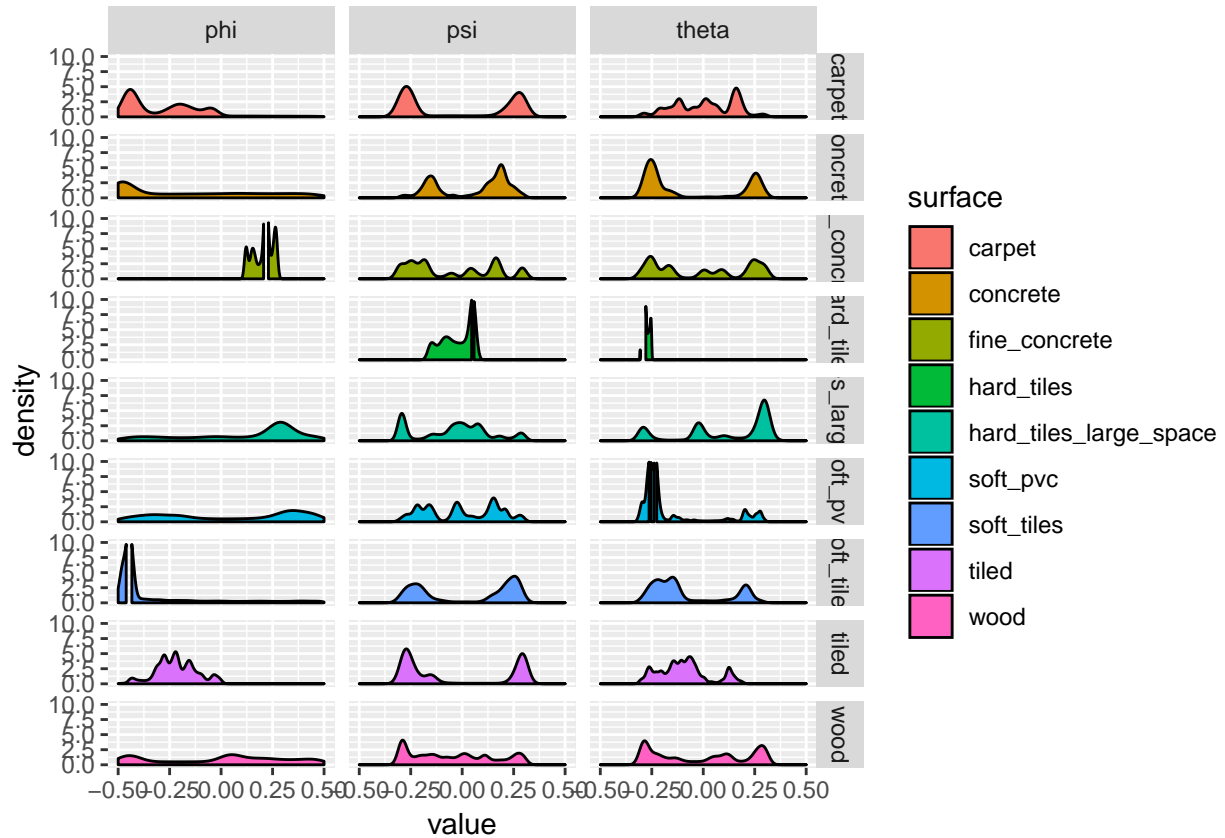
```

}

x_train <- convert_quaternions_to_euler(x_train)
x_test  <- convert_quaternions_to_euler(x_test)

```


Distribution of orientation angles by surface:



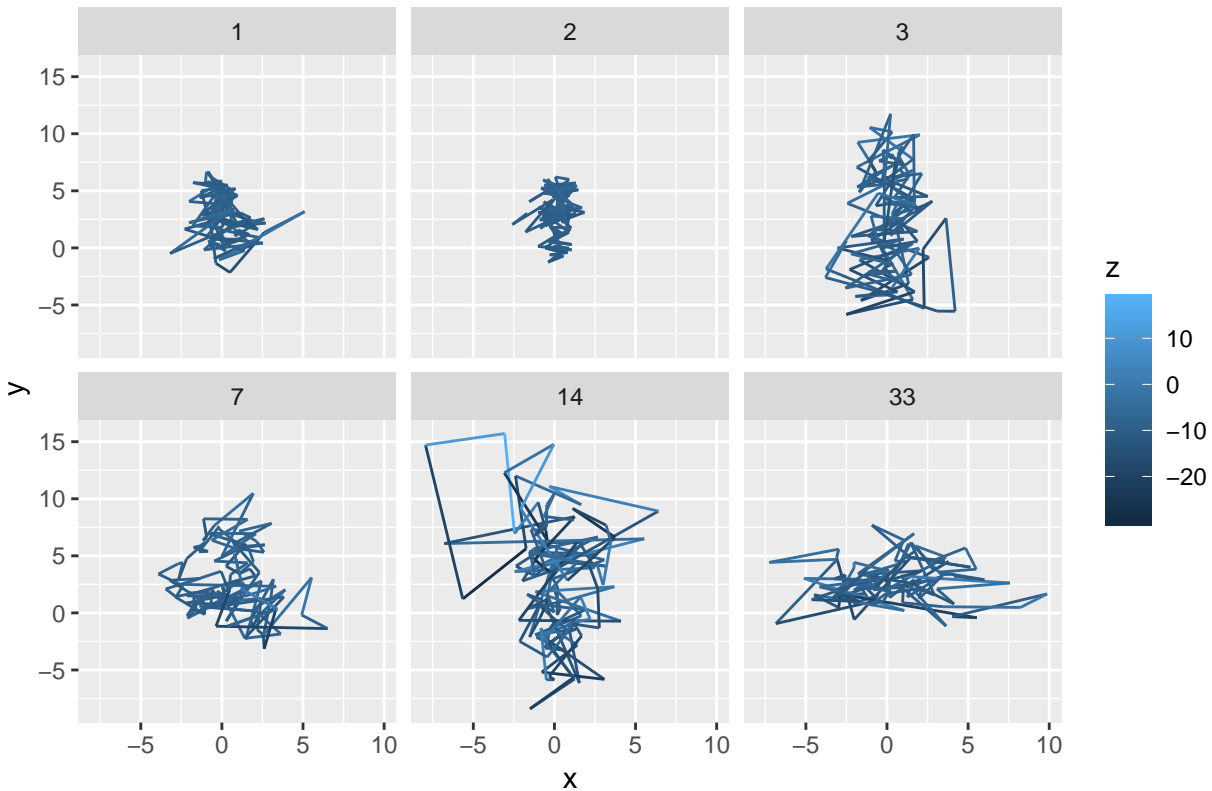
It looks like the most noticeable differences between surfaces can be observed in the orientation distribution.

If we consider unit of time to be 1, we can draw the path of the robot movement. In the following figures I draw the path for several cases, faceted by surface. I expect that the surface would influence the shape of these paths.

Robot path for surface = carpet



### Robot path for surface = concrete



I cannot observe obvious differences between paths on carpet vs concrete, but maybe a machine can. An idea would be to convert all paths to images and do an image analysis, but I will stick with a more classical approach for this project. I'll use several models in caret package.

## Data Pre-Processing

An important step in fitting a good model is pre-processing the data. In our case we have series of observation of 128 measurements. In this section I will create aggregations around these measurements. These are: mean, standard deviation, distances of segments from one point to the next. total distance, total angle change both from magnetostatic and gyro sensors. All of these are encapsulated in the following function:

```
# function for pre-processing
# n_of_rows defaults to total number of series
# is_train indicates that the data is training, thus will do an extra action
pre_process <- function(a_dataframe, n_of_rows = nrow(a_dataframe)/128 ) {

  # get data grouped by series_id and compute some means
  processed_data_df <- a_dataframe %>%
    group_by(series_id) %>%
    summarize(
      phi_mean_all = mean(phi),
      phi_sd_all = sd(phi),
      phi_mean_to_sd_all = mean(phi)/sd(phi),
      theta_mean_all = mean(theta),
```



```

theta_sd_all = sd(theta),
theta_mean_to_sd_all = mean(theta)/sd(theta),
psi_mean_all = mean(psi),
psi_sd_all = sd(psi),
psi_mean_to_sd_all = mean(psi)/sd(psi),
# this is the rectangular area that surrounds the path of linear movement
dist_area = (max(linear_acceleration_X) - min(linear_acceleration_X)) * (max(linear_acceleration_Z) - min(linear_acceleration_Z)) * (max(linear_acceleration_Y) - min(linear_acceleration_Y)) * (max(linear_acceleration_Z) - min(linear_acceleration_Z))
# this is the rectangular area that surrounds the path of angular movement
omega_area = (max(angular_velocity_X) - min(angular_velocity_X)) * (max(angular_velocity_Y) - min(angular_velocity_Y)) * (max(angular_velocity_Z) - min(angular_velocity_Z)) * (max(angular_velocity_Z) - min(angular_velocity_Z))
euler_area = (max(phi) - min(phi)) * (max(theta) - min(theta)) +
              (max(psi) - min(psi)) * (max(theta) - min(theta)) +
              (max(theta) - min(theta)) * (max(psi) - min(psi)),
dist_mean_x = mean(linear_acceleration_X),
dist_mean_y = mean(linear_acceleration_Y),
dist_mean_z = mean(linear_acceleration_Z),
omega_mean_x = mean(angular_velocity_X),
omega_mean_y = mean(angular_velocity_Y),
omega_mean_z = mean(angular_velocity_Z),
dist_sd_x = sd(linear_acceleration_X),
dist_sd_y = sd(linear_acceleration_Y),
dist_sd_z = sd(linear_acceleration_Z),
omega_sd_x = sd(angular_velocity_X),
omega_sd_y = sd(angular_velocity_Y),
omega_sd_z = sd(angular_velocity_Z)

) %>%
slice(1:n_of_rows)

# define an empty data frame with summary metrics that we'll use for each set of 128 observations
metrics <- c("dist_total", "dist_max", "dist_min", "dist_max_to_min", "dist_mean", "dist_sd", "dist_mean_to_sd",
             "omega_total", "omega_max", "omega_min", "omega_max_to_min", "omega_mean", "omega_sd", "omega_mean_to_sd",
             "phi_total", "phi_max", "phi_min", "phi_mean", "phi_sd", "phi_mean_to_sd",
             "theta_total", "theta_max", "theta_min", "theta_mean", "theta_sd", "theta_mean_to_sd",
             "psi_total", "psi_max", "psi_min", "psi_mean", "psi_sd", "psi_mean_to_sd",
             "euler_total", "euler_max", "euler_min", "euler_mean", "euler_sd", "euler_mean_to_sd")

tmp_df <- data.frame(matrix(ncol = length(metrics), nrow = 0) )
colnames(tmp_df) <- metrics

# loop over each series and compute aggregations
# should use apply type of function here, but I use "for" until I master the apply
for (s_id in processed_data_df$series_id)
{
  # get current measurement set
  this_chunk_df <- a_dataframe %>% filter(series_id == s_id)

  # select only columns we are interested in
  this_chunk_df <- this_chunk_df %>%
    select(
      phi, theta, psi,

```

```

    angular_velocity_X, angular_velocity_Y, angular_velocity_Z,
    linear_acceleration_X, linear_acceleration_Y, linear_acceleration_Z)
# this is a vector with euclidian distances from one point to the next
dist_v <- sqrt(diff(this_chunk_df$linear_acceleration_X)^2 +
               diff(this_chunk_df$linear_acceleration_Y)^2 +
               diff(this_chunk_df$linear_acceleration_Z)^2)
omega_v <- sqrt(diff(this_chunk_df$angular_velocity_X)^2 +
               diff(this_chunk_df$angular_velocity_Y)^2 +
               diff(this_chunk_df$angular_velocity_Z)^2)
phi_v <- abs(diff(this_chunk_df$phi))
theta_v <- abs(diff(this_chunk_df$theta))
psi_v <- abs(diff(this_chunk_df$psi))

# fill or temp data frame with summary computations for our 128 measurement set
tmp_df <- bind_rows(tmp_df, data_frame(

  # all features starting with "dist_" refers to linear movement
  dist_total = sum(dist_v),
  dist_max = max(dist_v),
  dist_min = min(dist_v),
  dist_max_to_min = max(dist_v)/min(dist_v),
  dist_mean = mean(dist_v),
  dist_sd = sd(dist_v),
  dist_mean_to_sd = mean(dist_v)/sd(dist_v), # reciprocal coef of variation

  # all features starting with "omega_" refers to angle velocity measurments
  omega_total = sum(omega_v),
  omega_max = max(omega_v),
  omega_min = min(omega_v),
  omega_max_to_min = max(omega_v)/min(omega_v),
  omega_mean = mean(omega_v),
  omega_sd = sd(omega_v),
  omega_mean_to_sd = mean(omega_v)/sd(omega_v), # reciprocal coef of variation

  # phi, theta and psi reffers to roll, pitch and yaw rotations
  phi_total = sum(phi_v),
  phi_max = max(phi_v),
  phi_min = min(phi_v),
  phi_mean = mean(phi_v),
  phi_sd = sd(phi_v),
  phi_mean_to_sd = mean(phi_v)/sd(phi_v),

  theta_total = sum(theta_v),
  theta_max = max(theta_v),
  theta_min = min(theta_v),
  theta_mean = mean(theta_v),
  theta_sd = sd(theta_v),
  theta_mean_to_sd = mean(theta_v)/sd(theta_v),

  psi_total = sum(psi_v),
  psi_max = max(psi_v),
  psi_min = min(psi_v),
  psi_mean = mean(psi_v),

```

```

psi_sd = sd(psi_v),
psi_mean_to_sd = mean(psi_v)/sd(psi_v),

# features starting with "euler_" refers to
# agragation of all phi, theta and psi rotations
euler_total = sum(phi_v + theta_v + psi_v),
euler_max = max(phi_v + theta_v + psi_v),
euler_min = min(phi_v + theta_v + psi_v),
euler_mean = mean(phi_v + theta_v + psi_v),
euler_sd = sd(phi_v + theta_v + psi_v),
euler_mean_to_sd = mean(phi_v + theta_v + psi_v)/sd(phi_v + theta_v + psi_v)
))

} # end of for over series

# add the summary computations to the data set of series
processed_data_df <- bind_cols(processed_data_df, tmp_df)

# more features
# I added these as an experimentation after observing the dependencies graphs
# will explain later
processed_data_df <- processed_data_df %>% mutate(
  # this is an agular momentum
  f1 = log(dist_mean_to_sd*omega_mean_to_sd),

  # this is an angular momentum
  f2 = log(dist_total*omega_total),

  # this is the angle between theta and psi vectors
  f3 = abs(atan(theta_mean_all/psi_mean_all)))

# return the proceessed data
processed_data_df
}

```

Pre-process the data and save it to a local file. This is to save time for further analysis by skipping the pre-processing procedure. I use `tic()` `toc()` functions from `tictoc` package to display processing time.

```

# pre-process train and test data sets
# tic("process train data")
# x_train_processed <- pre_process(x_train)
# toc()
#
# tic("process test data")
# x_test_processed <- pre_process(x_test)
# toc()
#
# # rejoin train data with the labels data set
# x_train_processed <- x_train_processed %>% left_join(y_train, by = "series_id")
#
# # creeatee a folder "data" if doesnt exist
# if (!dir.exists("data")) dir.create("data")
#

```

```

# # save processed data, and use these files from now on
# write_csv(x_train_processed, "data/x_train_processed.csv")
# write_csv(x_test_processed, "data/x_test_processed.csv")
#
# # clean some variables and the environment
# rm(x_train, x_test, y_train)
# rm(x_train_processed, x_test_processed)

```

Load pre-processed data from hard-disk and use them from here. In this report I will use a smaller data set from the training data, to make things work faster. Full data set will be used in the final script

```

x_train_processed_from_file <- read_csv("data/x_train_processed.csv")
x_test_processed_from_file <- read_csv("data/x_test_processed.csv")

# if we load data from a file, convert surface to factor
x_train_processed_from_file <- x_train_processed_from_file %>% mutate(surface = as.factor(surface))

# use a smaller set of datab to save time
x_train_processed_from_file <- x_train_processed_from_file %>% slice(1:1200)

```

## Correlated Predictors

```

# from now we can load pre-processed data from csv files without the need to pre-process again
x_test_processed <- x_test_processed_from_file
x_train_processed <- x_train_processed_from_file

```


We created 65 features, lets trim some of them. I will identify first the ones that are strongly correlated.

```

library(RColorBrewer)
library(corrplot)
x_matrix <- x_train_processed %>% select(-series_id, -group_id, -surface) %>%
  as.matrix()
x_cor <- cor(x_matrix, use = "pairwise.complete")
corrplot(x_cor, method = "square", number.cex = .5, tl.cex = 0.6, order = "hclust", title = "Correlation")

```

## Correlation Matrix



```
rm(x_matrix, x_cor)
```

Using function *findCorrelation* from *caret* we can remove features that are highly correlated

```
# convert both test and train data to matrix in order to analyse feature correlation
x_train_matrix <- x_train_processed %>% select(-surface, -series_id) %>% as.matrix()
x_test_matrix <- x_test_processed %>% select(-series_id) %>% as.matrix()

# find features that are high correlated
```

```


# find linear dependencies and eliminate them
names_to_remove_test <- findCorrelation(cor(x_test_matrix), cutoff = 0.95, names = TRUE, verbose = FALSE)

# remove correlated features from both train and test sets
x_train_processed <- x_train_processed %>% select(-names_to_remove_test)
x_test_processed <- x_test_processed %>% select(-names_to_remove_test)

# draw again the correlation matrix
x_matrix <- x_train_processed %>% select(-series_id, -group_id, -surface) %>%
  as.matrix()
x_cor <- cor(x_matrix, use = "pairwise.complete")
corrplot(x_cor, method = "square", number.cex = .5, tl.cex = 0.6, order = "hclust", title = "Correlation")

```

## Correlation Matrix



```
rm(x_train_matrix, x_test_matrix, x_matrix, x_cor)
```

The number of predictors decreased to 42

Another useful technique in pre-processing is centering around means and scaling the predictors:

```
# pre-process the data, center and scale the values across all predictors
pre_process <- x_train_processed %>% select(-series_id, -group_id) %>% preProcess(method = c("center",
x_train_processed <- predict(pre_process, x_train_processed)
```


```
x_test_processed <- predict(pre_process, x_test_processed)

rm(pre_process)
```

## Visualizations


Before moving forward with more feature selection, I will plot some interesting visualizations of the feature dependencies. These are plots of orientation angles as sums over all 128 points.  $\phi$ ,  $\theta$  and  $\psi$  are euler orientation angles.

```
x_train_processed %>% ggplot(aes(phi_mean_all, psi_mean_all, fill = surface)) +
  geom_point(aes(color = surface)) +
  geom_point(cex=6, pch=21) +
  ggtitle("phi_mean vs. psi_mean")
```




```
# this is a circle
x_train_processed %>% ggplot(aes(theta_mean_all, psi_mean_all, fill = surface)) +
  geom_point(aes(color = surface)) +
  geom_point(cex=6, pch=21) +
  ggtitle("theta_mean vs. psi_mean")
```






```
# this is a wave
x_train_processed %>% ggplot(aes(theta_mean_all, phi_mean_all, fill = surface)) +
  geom_point(aes(color = surface)) +
  geom_point(cex=6, pch=21) +
  ggtitle("theta_mean vs. phi_mean")
```



I would speculate that all points that are not on the global pattern are actually outliers, but I did not explore that hypothesis.

The following is the angle between  $\theta$  -  $\psi$  vectors and the angular momentum:

```
x_train_processed %>% ggplot(aes(f2, f3, fill = surface)) +  
  geom_point(aes(color = surface)) +  
  geom_point(cex=6, pch=21) +  
  ggtitle("theta-psi angle vs. angular momentum")
```



We can see some separations by surfaces that may allow a KNN model to work. In fact if we run KNN for just  $\text{surface} \sim f2 + f3$  we will get an accuracy of 0.56 which is promising.

## More Feature Selection

I will continue now removing more features based on their importance in a randomForest model. I will split the training set into two partitions:  $x_{\text{train\_for\_train}}$  used for training, and  $x_{\text{train\_for\_test}}$  to measure the accuracy. Then will plot variable importance. Will not use here any tuning, since this is not the final model.

```
# partition x_train_processed data for training and testing  
test_index <- createDataPartition(y = x_train_processed$surface, times = 1, p = 0.5, list = FALSE)  
x_train_for_train <- x_train_processed[-test_index, ]  
x_train_for_test <- x_train_processed[test_index, ]  
rm(test_index)  
  
model_fit <- randomForest(  
  surface ~ .,  
  metric = "Accuracy",  
  # remove series_id, group_id and select only 500 points
```

```

data = slice(select(x_train_for_train, -series_id, -group_id), 1:500)
)


y_hat <- predict(model_fit, select(x_train_for_test, -series_id))
y_test <- x_train_for_test$surface

# # we dont show accuracy and confusion table for now
# conf_matrix <- confusionMatrix(y_hat, x_train_for_test$surface)
# conf_matrix$overall["Accuracy"]
# conf_matrix$table %>% knitr::kable()

importance <- importance(model_fit)
importance <- importance[order(importance[,1], decreasing = TRUE), ]

# plot variable importance
varImpPlot(model_fit)

```



```

# this is the list of variables in order of importance, decreasing
features <- names(importance)

```

We observe following variables does not help prediction too much: theta\_min, omega\_max\_to\_min, dist\_mean\_y, omega\_mean\_x, dist\_mean\_x, dist\_mean\_z So I will remove them:

```

# remove columns do not contribute to classification
x_train_processed <- x_train_processed %>% select(-theta_min, -omega_max_to_min, -dist_mean_y, -omega_mean_x, -dist_mean_x, -dist_mean_z)
x_test_processed <- x_test_processed %>% select(-theta_min, -omega_max_to_min, -dist_mean_y, -omega_mean_x, -dist_mean_x, -dist_mean_z)

```

## Find the Best Model

It is time now to fit some models. I tried many models on the side but I will present in this report only

- KNN,
- rpart,
- randomForest
- randomForest assembly

Please note that the code in this report is using only 1000 records from training, that is further partitioned in train and test with 0.5 ratio I've done this to make it run faster. I am using full data set in the R script and also I will present the final results from Kaggle.


## KNN

KNN is getting incrementally slow with the number of features. So for KNN I used only the most important variables as part of the formula.

```
# create a data frame to store Accuracy results by model
model_results <- data.frame(method = "Model", accuracy = "Accuracy")

model_fit <- train(surface ~ f3 + f2 + phi_mean_all + theta_mean_all + psi_mean_all, method = "knn",
  tuneGrid = data.frame(k = seq(2, 100, 2)),
  data = select(x_train_for_train, -series_id, -group_id))

ggplot(model_fit)
```



```
# get confusion matrix and display it together with the results
y_hat <- predict(model_fit, x_train_for_test, type = "raw")
conf_matrix <- confusionMatrix(y_hat, x_train_for_test$surface)

# display confusion matrix
conf_matrix$table %>% knitr::kable()
```

	carpet	concrete	fine_concrete	hard_tiles	hard_tiles_large_space	soft_pvc	soft_tiles
carpet	20	4	0	0	0	5	0
concrete	5	84	0	0	0	5	0
fine_concrete	0	0	33	0	5	1	1
hard_tiles	1	0	0	7	0	0	0
hard_tiles_large_space	0	2	0	0	23	2	0
soft_pvc	4	4	4	0	2	114	7
soft_tiles	1	1	4	0	6	8	91
tiled	7	5	0	1	0	0	2
wood	0	2	2	0	1	2	1

```
model_results <- bind_rows(model_results,
  data.frame(
    method = "KNN",
    accuracy = as.character(conf_matrix$overall["Accuracy"])))
```

```
model_results %>% knitr::kable()
```

method	accuracy
Model	Accuracy
KNN	0.764510779436153

## One-vs-One or One-vs-All

## Results and Submit the Data

## Conclusion

One of the most important outcome of this prproject is that I learned a few things in addition to what was presented in the course.

## Reference

1. Applied Predictive Modeling - Max Kuhn, Kjell Johnson
2. Q2EA: Convert from rotation Quaternions to Euler Angles. Q2EA converts from Quaternions (Q) to Euler Angles (EA) based on D. M. Henderson (1977). Q2EA.Xiao is the algorithm by J. Xiao (2013) for the Princeton Vision Toolkit. <https://rdrr.io/cran/RSpincalc/man/Q2EA.html>
3. Understanding Quaternions. <http://www.chrobotics.com/library/understanding-quaternions>
4. Understanding Euler Angles. <http://www.chrobotics.com/library/understanding-euler-angles>
5. Tune Machine Learning Algorithms in R (random forest case study) by Jason Brownlee. <https://machinelearningmastery.com/tune-machine-learning-algorithms-in-r/>
6. Classification with more than two classes, from Introduction to Information Retrieval, Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze,, Cambridge University Press 2008 <https://nlp.stanford.edu/IR-book/html/htmledition/classification-with-more-than-two-classes-1.html>
7. A Guide To using IMU (Accelerometer and Gyroscope Devices) in Embedded Applications - [http://www.starlino.com/imu\\_guide.html](http://www.starlino.com/imu_guide.html)